

Executive summary

Monitored target: Main RGDA

Coverage start: Jul 31, 2026 at 5:18 PM PDT

Coverage end: Aug 3, 2026 at 4:44 PM PDT

Report window: 3.0 days (inaugural partial week)

31.5 h Scheduled monitoring time	\$32.80 Projected API-equivalent cost	26 Incidents detected	26 Incidents resolved / closed	0 High / critical unresolved	8 Configured role threads
--	---	---------------------------------	--	--	-------------------------------------

EFFECTIVE-WITH-FINDINGS

All 26 incidents reached terminal outcomes; recurrence and validator coupling remain the key risks

Executive supervisor assessment

- Detection: Supervisor replay-cost review eliminated unnecessary deep work
- Effectiveness: Every incident reached a terminal outcome by the cutoff
- Residual risk: Mechanical supervisor validation can share a producer blind spot
- Efficiency: The monitor found expensive waste, often after one extra attempt

What was running

Routine watcher — 28 watcher checks Mechanical change gate and scheduled target checks. (configured)	Semantic reviewer — 156 semantic reviews Substantive review of each routed changed target state. (configured)
Max reviewer — 38 sample/effectiveness reviews Sampled changed-state review plus scheduled effectiveness review. (configured)	Incident outcome reviewer — 23 incident review actions Material incident adjudication and correction-outcome review. (configured)
Fix executor — 10 fix actions Independently approved, bounded fixes to supervision machinery. (configured)	Roundup/report writer — 11 roundups Scheduled summaries and cognitive weekly supervision review. (configured)
Gmail reply gate — 0 reply-gate records Validates whether an inbound roundup-thread reply may be processed. (configured)	Gmail reply processor — 0 reply-processing records Processes only replies admitted by the maintained Gmail gate. (configured)

Configured threads describe the maintained supervision roles. Recorded actions are ledger-visible lower bounds because unchanged scheduled wakes may intentionally emit no record.

Inside this report


Monitoring time, read reliability, and projected cost
Supervisor effectiveness
Coverage and operating efficiency
Methodology and evidence boundary

Activity, incidents, and response time
Detection quality
Monitoring machinery evolution

Monitoring time, read reliability, and projected cost


This page separates explicit schedule posture, recorded target-read reliability, and projected resource use. It does not infer continuous process uptime from quiet ledger gaps.

Total report period	71.42 h
Scheduled monitoring time	31.49 h (44.09%)
Explicitly paused monitoring time	39.93 h
Recorded target-read reliability	92.39% (170 successful / 14 failed reads)
Continuous process uptime measured	No



Scheduled-active time is derived only from explicit pause/resume policy records. Target-read availability uses recorded material read outcomes. Quiet gaps are not downtime because unchanged wakes may emit no event.

Estimated token use and projected API-equivalent cost




Projected API-equivalent cost (USD)				
Low estimate — Lower token multiplier from the versioned estimation profile.		Base estimate — Central API-equivalent projection used in the executive summary.		High estimate — Upper token multiplier; still a projection, not billed usage.
Model	Records	Estimated tokens	Projected cost	API pricing assumption
gpt-5.6-sol	685	5,756,490 / 3,453,894-9,210,384	\$31.91 / \$19.15-\$51.06	GPT-5.4 standard API equivalent
gpt-5.6-terra	170	461,983 / 277,190-739,173	\$0.88 / \$0.53-\$1.41	GPT-5.4 mini standard API equivalent
TOTAL	855	6,218,473 / 3,731,084-9,949,557	\$32.80 / \$19.68-\$52.47	supervision-weekly-cost-projection-v1


Activity, incidents, and response time

These charts show ledger-visible supervision activity by local day. Their units and categories describe the monitoring system, not implementation quality.

Recorded monitoring activity by day



Incident detection and resolution by day



Incident response posture

Detection yield	15.29 incidents per 100 routed changed states
Resolved / closed share	100.0% of incidents detected in the period
Median detection-to-resolution time	0.71 h
P90 detection-to-resolution time	4.97 h
Unresolved incidents at cutoff	0 total; 0 high/critical

Incident rate uses incident openings divided by exact changed-state routing records. It measures supervision yield, not implementation quality.

Supervisor effectiveness

Detection value and correction effectiveness

What supervision caught or prevented

Supervisor replay-cost review eliminated unnecessary deep work

Supervision detected an ordinary replay path that was rebuilding a 2,134-file evidence envelope despite a current frozen baseline. The correction moved the normal monitoring proof to cheap revision and root checks while retaining a deep check when currentness actually changes. This improved execution economy without weakening fail-closed detection.

Evidence: EVT-000149, EVT-000162

Independent semantic review caught circular validation

Mechanical checks initially agreed because validation reproduced the producer classification logic. Independent semantic review reconstructed the governing facts, kept promotion closed, and later verified the bounded successor. This demonstrates the value of layered supervision and the risk of mechanically coupled assurance.

Evidence: EVT-000886, EVT-000896, EVT-000904, EVT-000919

Policy-hash checks exposed stale role activation

A current-policy hash mismatch prevented a wrong review and exposed stale automation and role prompts. The supervisor repaired only those prompt references, preserved schedules and identities, and verified the current contract before closing the incident.

Evidence: EVT-000925, EVT-000926, EVT-000932, EVT-000933

Whether supervisor corrections worked

Every incident reached a terminal outcome by the cutoff

All 26 opened incidents reached corrected, resolved, or closed posture by the cutoff. Median detection-to-resolution time was 0.71 hours and the 90th percentile was 4.97 hours. This shows responsive incident handling; recurrence still prevents treating terminal closure as proof that each monitoring lesson is durable.

Evidence: EVT-000031, EVT-000049, EVT-000919, EVT-000933

Command-envelope monitoring became durable

An early runtime correction did not bind every repository import and selection assumption, so the setup failure recurred. Supervision then moved the complete maintained command envelope ahead of invocation. Later applicable audits collected and ran on their first attempt, providing direct effectiveness evidence.

Evidence: EVT-000177, EVT-000294, EVT-000770, EVT-000771

Mission-provenance review repaired a false workflow stop

A narrow historical no-rerun instruction was incorrectly treated as a durable ban. Supervisor review preserved the mistaken history, corrected the authority interpretation, resumed ordinary successor work, and added a reviewed provenance rule. This fixed the monitoring decision path rather than converting process preservation into a new user prohibition.

Evidence: EVT-000730, EVT-000732, EVT-000746, EVT-000769

Detection quality

Recurrence, false positives, misses, and blind spots

Recurring supervision failure classes

Supervision repeatedly encountered known lower-path ordering failures

Several incidents began with a broad or incomplete operation before using an already-proven isolated path and exact exclusions. Individual incidents closed, but recurrence shows that terminal closure alone did not make the correction durable; prebound ordering was the more effective monitoring control.

Evidence: EVT-000326, EVT-000363, EVT-000513, EVT-000649

Execution preflight often occurred after expensive work began

Runtime, import, reporting, packet-bound, and capacity checks were sometimes discovered only after costly work. Supervisor maintenance moved those checks before execution and preserved valid producer evidence when only reporting failed. This generalized pattern explains several avoidable retries.

Evidence: EVT-000294, EVT-000617, EVT-000680, EVT-000686

Historical holds repeatedly drifted into current authority claims

Across several monitoring decisions, operation-scoped or expired language was treated as if it governed later successor work. Mission and containment provenance checks reduced this class, but later recurrence shows that supervisors must verify scope and expiry at the decision point.

Evidence: EVT-000730, EVT-000836, EVT-000895, EVT-000901

False positives, misses, and blind spots

Mechanical supervisor validation can share a producer blind spot

One consequential semantic result passed producer and validator logic because both used the same classifier. Independent review later found and closed the defect. Consequential semantic dispositions therefore need independently reconstructed governing facts, not agreement between coupled implementations.

Evidence: EVT-000886, EVT-000896, EVT-000908, EVT-000919

Some supervisor catches occurred after a live effect

Two material incidents were detected only after a nonfinal live operation had occurred. Recovery preserved history and enabled correction, but these episodes show that monitoring performance was not purely preventive and that pre-promotion review timing remains important.

Evidence: EVT-000568, EVT-000609, EVT-000895, EVT-000896

No formal false-positive closure is not proof of zero oversteer

The metrics contain no incident terminated with the formal false-positive status, yet supervision created one false reserved-authority block and issued some corrections that later proved overly broad or incomplete. Formal status counts therefore understate decision-quality errors unless paired with cognitive review.

Evidence: EVT-000480, EVT-000503, EVT-000730, EVT-000769

Coverage and operating efficiency

Pace, resource posture, and avoidable supervision cost

Coverage and monitored development pace

Monitoring covered rapid movement across 29 tracker stages

Supervision observed 29 distinct tracker stages during 71.33 hours, with repeated checkpoint and transition evidence. This demonstrates broad monitoring coverage and sustained implementation movement; it does not mean every observed stage was accepted or substantively correct.

Evidence: EVT-000436, EVT-000452, EVT-000877, EVT-000922

Supervisor corrections generally preserved forward motion

The most effective interventions froze valid checkpoints, retained passing evidence, recomputed only affected proof, and kept dependent work closed without stopping independent safe work. This reduced the delivery cost of correction relative to replaying whole completed slices.

Evidence: EVT-000774, EVT-000858, EVT-000874

High supervisor intensity accompanied the monitored pace

The window contains 933 records, 169 changed-state routes, 156 semantic reviews, 17 Max samples, 21 effectiveness reviews, and 97 email receipts. The coverage is substantial but operationally expensive; recorded activity is a lower bound and not an automatic productivity metric.

Evidence: EVT-000875, EVT-000881, EVT-000924, EVT-000933

Supervisor operating efficiency

The monitor found expensive waste, often after one extra attempt

Supervisor catches included a full 2,134-file replay, repeated large-database hashing, broad-operation rollbacks, malformed packet retries, and reporting-only failures that repeated valid proof. Corrections were narrow and reusable, but the pattern shows a recurring one-extra-attempt cost before the proven lower path was selected.

Evidence: EVT-000149, EVT-000308, EVT-000649, EVT-000678, EVT-000853

Supervisor activity counts require cadence context

Only 28 watcher checks are recorded because unchanged scheduled wakes may emit no event, while every distinct changed state is expected to route explicitly. The report therefore compares intended cadence with recorded activity instead of presenting event counts as exact scheduler-run counts.

Evidence: EVT-000875, EVT-000878, EVT-000881

Monitoring time, read reliability, and model economics have explicit boundaries

The report covers 71.33 hours: 31.40 scheduled-active monitoring hours and 39.93 explicitly paused hours. Recorded target-read reliability was 92.35 percent. The profile projects 6.21 million base tokens and 32.75 dollars of API-equivalent cost. These are not continuous uptime, provider telemetry, actual internal prices, or billed cost.

Evidence: EVT-000220, EVT-000754, EVT-000933

Monitoring machinery evolution

What changed in the supervisor and what should improve next

Monitoring machinery change log

Notification and roundup monitoring became closed-loop

The supervision system added quiet-window Gmail gating, event-driven incident review, source-linked terminal outcomes, scheduled roundups, lifecycle delivery, and deterministic notification deduplication. A later helper fix prevented delivery receipts from replacing substantive incident heads.

Evidence: EVT-000027, EVT-000039, EVT-000083, EVT-000562

Mission and authority provenance became explicit in supervision

A reviewed maintenance sequence bound the monitor to direct mission sources and challenged supervisor-created reserved authority, expired containments, and goal-reversing holds. The improvement corrected a real obstruction, although later recurrence shows the check must run before consequential intervention.

Evidence: EVT-000740, EVT-000746, EVT-000749, EVT-000895

Weekly reporting and role activation became reproducible

The supervisor now produces deterministic metrics, hash-bound cognitive review, canonical JSON, Markdown, charted PDF, verification, and same-thread delivery. It also detected and repaired stale role prompt references while preserving schedules, identities, and other automation fields.

Evidence: EVT-000912, EVT-000917, EVT-000927, EVT-000932, EVT-000933

Improvements to supervision machinery

Require independent semantic reconstruction in supervisor validation

For consequential semantic dispositions, monitoring should compare the producer result with an independently reconstructed set of governing facts. A validator that merely repeats producer classification logic should be marked mechanically coupled and escalated before the supervisor treats it as acceptance evidence.

Evidence: EVT-000886, EVT-000896, EVT-000904, EVT-000908

Measure supervision recurrence separately from incident closure

Weekly monitoring should retain both terminal closure rate and recurrence by generalized failure class. Ordering, command-envelope, and authority-provenance defects recurred after earlier terminal incidents, so closure percentages without recurrence context overstate supervisor effectiveness.

Evidence: EVT-000326, EVT-000363, EVT-000294, EVT-000705, EVT-000730, EVT-000901

Verify supervisor role prompt activation at policy change

After a supervision policy change, the existing activation path should compare every configured role and automation prompt with the accepted policy and skill hashes before the next substantive wake. The check should repair only stale references and preserve schedules, identities, and other fields.

Evidence: EVT-000925, EVT-000926, EVT-000932, EVT-000933

Methodology and evidence boundary

How to interpret this supervisor-performance report

Methodology and limits

Every bounded supervision record entered deterministic computation

The canonical packet covers every validated event from EVT-000001 through EVT-000933 and all 23 in-window policy records. Metrics, incident heads, activity, schedule posture, role counts, and resource estimates were computed before cognitive interpretation.

Evidence: EVT-000001, EVT-000933

Supervisor synthesis is evidence-bound but not causal proof

The review explains monitoring patterns that raw counts cannot, but it cannot establish why target quality changed, whether unlogged no-op wakes occurred, continuous process uptime, exact provider tokens or billed cost, or whether any monitored substantive result is legally or technically adequate.

Evidence: EVT-000053, EVT-000879, EVT-000916

This supervisor report is an inaugural partial-week baseline

The available history spans 71.42 hours. Week-over-week trends in monitoring detection, recurrence, pace, read reliability, and cost require later reports using the same definitions, pricing profile, and evidence-bound cognitive-review contract.

Evidence: EVT-000001, EVT-000879, EVT-000933

Exact evidence binding

Report ID: weekly-20260801T001853Z-20260803T234400Z-62f881d083b0

Source root: 62f881d083b0e64262f54a81c575d0dd7b7a3e4330ea363f5a2b9414aadfd032

Source records: EVT-000001 through EVT-000935 (935 records)

The PDF is a deterministic supervisor-performance projection. It does not direct the monitored implementation, confer tracker completion, or establish patent or legal quality.

Resource projection boundary

For each model-attributed supervision record: canonical visible record characters divided by the configured character-per-token ratio, plus a model-specific input-context allowance; visible output fields or a model-specific output floor, multiplied by the recorded reasoning-effort factor. Low/high bounds multiply the base estimate. Costs apply the stated public API-equivalent price assumptions.

This profile projects API-equivalent cost from content-minimized supervision records. It is not provider token telemetry, actual usage, billed cost, an invoice, or the price of the internal supervision model aliases.

- Recorded task activity is a lower bound when scheduled no-op wakes intentionally emit no event.
- Counts and hashes establish operational completeness, not patent quality or legal sufficiency.
- A higher incident rate may reflect better detection or a riskier implementation phase; cognitive review is required before interpretation.
- The first report may cover less than seven days when supervision history is newer than one week.
- Token counts and costs are projection ranges from a versioned estimation profile, not provider telemetry, actual usage, billed cost, or the prices of internal model aliases.
- Continuous process uptime is not available from the content-minimized event ledger; only explicit scheduled-active/paused intervals and recorded target-read outcomes are reported.